

When Smaller Is Slower: Dimensional Collapse in Compressed LLMs

Suggested alternatives: **ShapeGuard: Preventing Dimensional Collapse in Compressed LLMs**; **GAC: GPU-Aligned Compression for Fast LLM Inference**.

Motivation

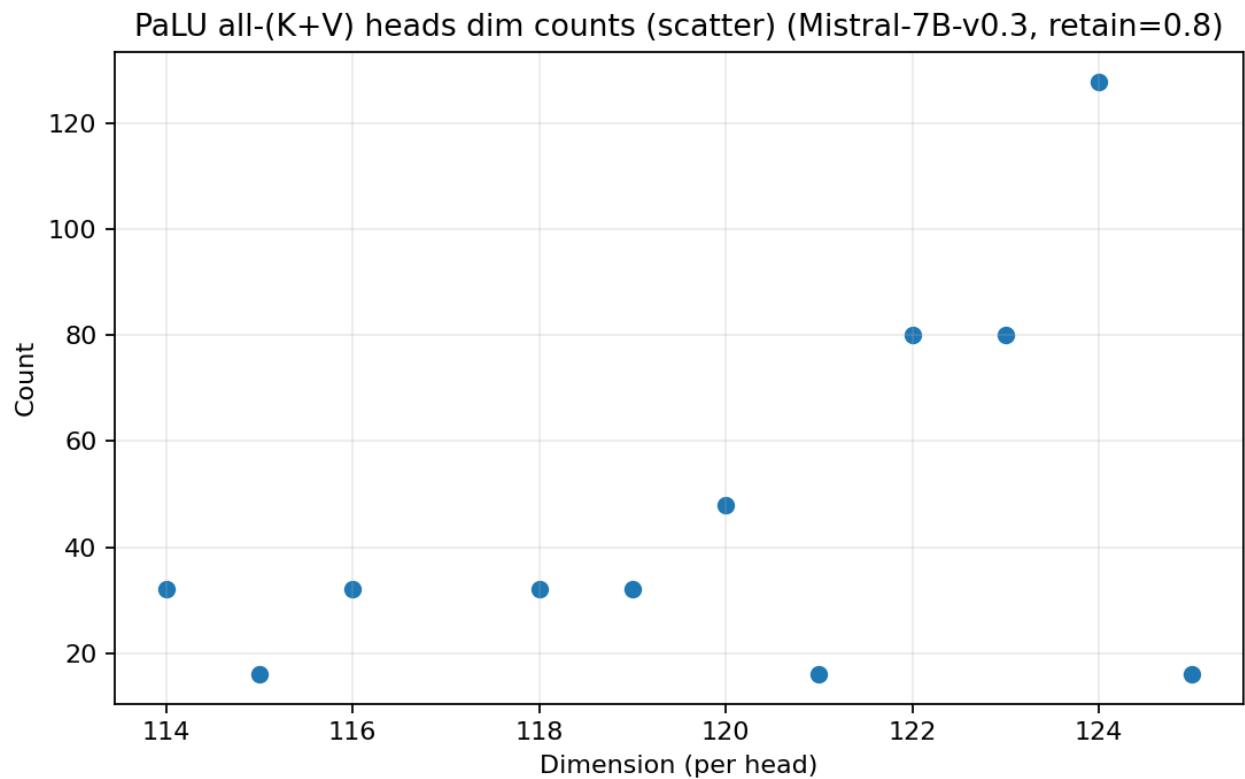
Post-training compression reshapes LLM operators (e.g., per-head dimensions become heterogeneous and irregular).

Irregular dimensions can trigger nonlinear GPU slowdowns (“dimensional collapse”), even when FLOPs decrease.

Kernel-aligned constraints (pad/pack/round) can rescue performance with small memory overhead.

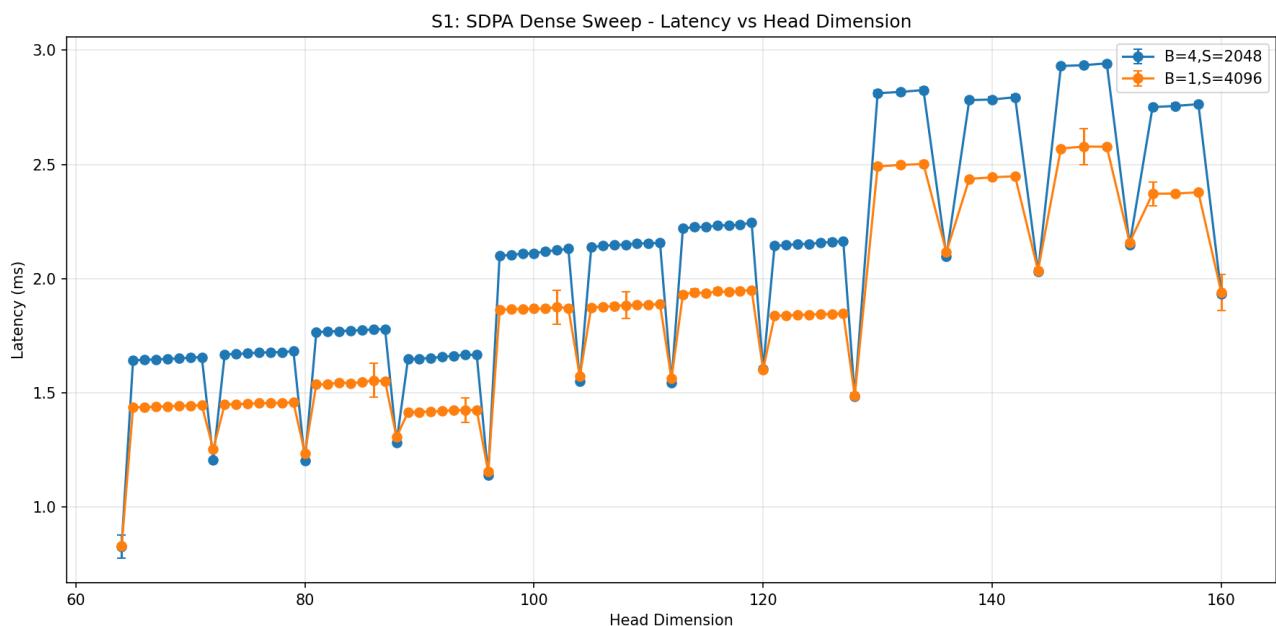
1. Dimensional Collapse

1.1 Compression changes per-head dimensions



Example: PaLU dimension distribution for Mistral (retain=0.8); per-head K/V dims vary.

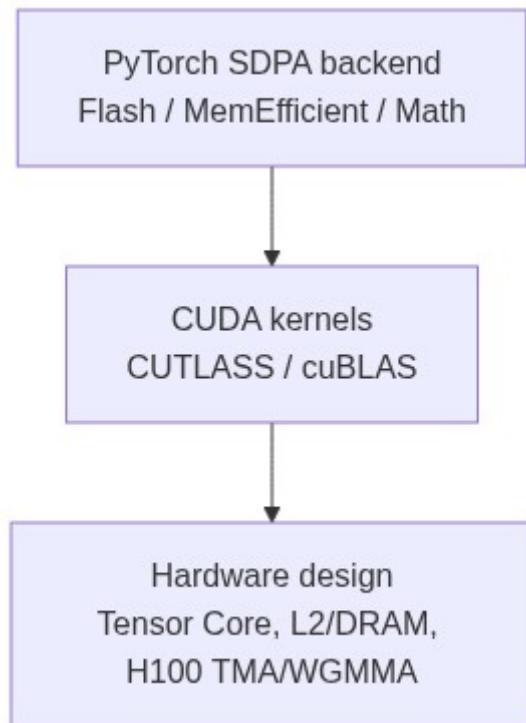
1.2 SDPA latency shows alignment cliffs



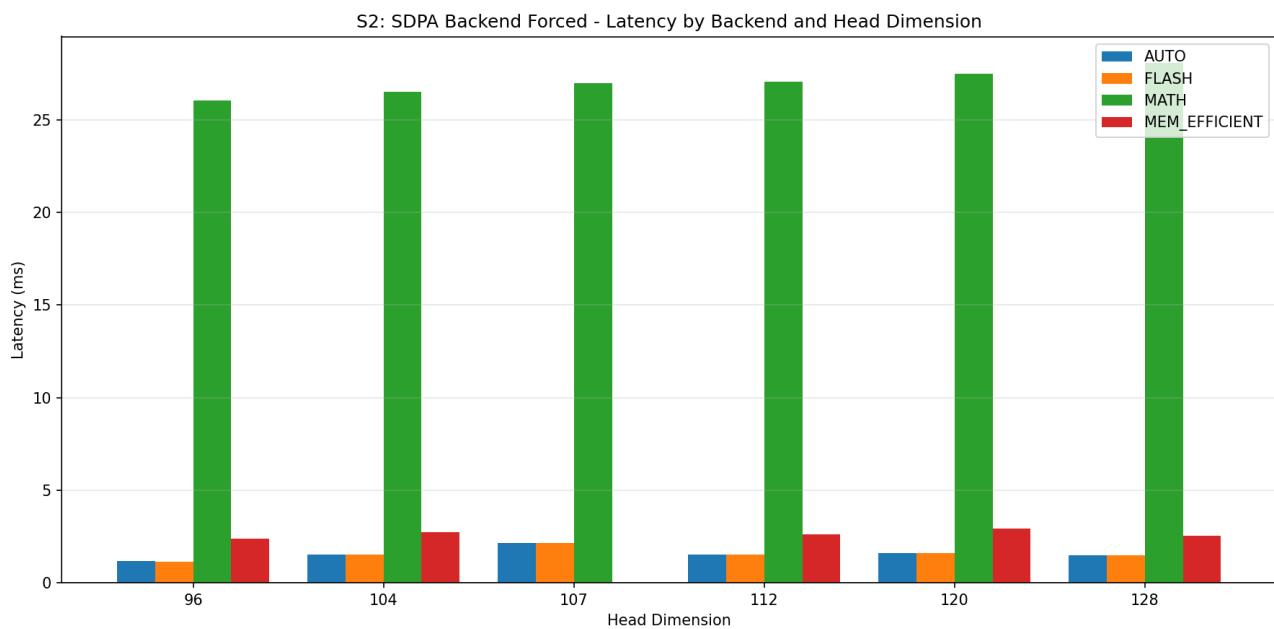
Experiment config (S1 / A100): B=4, S=2048, H=32 , dtype= fp16 ,
is_causal=False , warmup=50, measure=200, trials=3.

Key points: D=96: 1.140 ms , D=107: 2.147 ms (+88% vs 96), D=112: 1.545 ms (~-28% vs 107).

2. Possible Cause (Hypotheses)



2.1 PyTorch backend selection (evidence)



Experiment config (S2 / A100): `B=4, S=2048, H=32, dtype= fp16, is_causal=False, head_dim ∈ {96,104,107,112,120,128}.`

Observed at D=107 : `FLASH≈2.139 ms, MATH≈26.995 ms (~12.6x); AUTO≈FLASH ; MEM_EFFICIENT unsupported in this run.`

2.2 CUDA kernel layer (hypotheses)

- Kernel variants are heavily tiled/vectorized; irregular `D` can force predication, worse memory coalescing, or less-optimized kernels.
- CUTLASS/cuBLAS GEMM often prefers aligned `K/N` (Tensor Core tiles); irregular dims may trigger slower paths or extra packing.

2.3 Hardware layer (hypotheses)

- Tensor Cores prefer tile-compatible dimensions (often `K % 16 == 0` for FP16/BF16).
- L2/DRAM transaction granularity makes misaligned row strides waste bandwidth.
- On H100, TMA/WGMMA introduce stricter alignment/stride contracts; irregular strides can disable the fastest data-movement/compute paths.

EuroMLSys

We frame this project as a **new post-training compression paradigm**: compression must satisfy a GPU “shape contract”, otherwise irregular operator shapes can trigger dimensional collapse.

C1) Quantify the phenomenon

C2) How to probe the system

C3) Formulate as a constraint optimization problem

C4) Solver + repair pass (optional)

C5) Kernel Level Validation